

[⟨ Back to jobs](#)

New

Senior Human-AI Collaboration Research Engineer, Education Labs

[Apply](#)

San Francisco, CA

About Anthropic

Anthropic's mission is to create reliable, interpretable, and steerable AI systems. We want AI to be safe and beneficial for our users and for society as a whole. Our team is a quickly growing group of committed researchers, engineers, policy experts, and business leaders working together to build beneficial AI systems.

About the role

Education Labs is a research-driven team with a product mandate: we study how AI transforms human capability development, and we build new product paradigms based on what we discover. As a Senior Product Research Engineer, you'll be the second technical builder on this small, mission-driven team—operating with research-level autonomy while shipping features that reach millions of users.

This role sits at the intersection of research exploration and product execution. You'll work alongside researchers studying the societal impacts of AI, while building and shipping features that help people develop skills—both in working with AI itself and in their broader capabilities. You'll have significant creative license to define what “good” looks like, exploring new interaction paradigms rather than optimizing existing patterns.

Our team believes that developing skill with AI is fundamental to human agency and flourishing in an AI-enabled world. We're building experiences that make users more capable, more thoughtful, and more empowered—not just more engaged. This means

[Privacy](#) • [Terms](#)

integrating skill development principles into product design, leveraging Claude itself as part of the capability-building journey, and measuring success by how users grow in their ability to work with AI and apply those skills to their own goals.

You'll be something of a one-person technical shop with full-stack ownership, prototyping new ideas, establishing technical direction, and shipping production-quality features. You'll need strong product instincts and clean interface design sensibilities, balanced with comfort in ambiguity and frontier thinking.

Responsibilities:

- **Build and ship features** that help users develop skills in working with AI and translate those skills into their own capability growth—measuring success by skill development and product mastery, not just engagement
- **Architect and implement end-to-end prototypes** (front-end and back-end) that test new interaction paradigms for skill development with AI, including novel uses of Claude as a skill-building partner
- **Define technical direction and quality standards** for capability-focused product work in a research context—you'll establish patterns others can follow
- **Collaborate with researchers** to translate insights about skill development and human-AI collaboration into shipped product experiences
- **Work closely with product, design, and research teams** across Anthropic to influence how skill development principles shape our broader product strategy
- **Analyze product data and user behavior** using quantitative and qualitative methods to understand how users develop skills with AI and what actually helps them become more capable
- **Document and share your work** through clear technical writing, prototypes, and presentations that influence thinking across the organization
- **Elevate team craft** through thoughtful code reviews, technical discussions, and knowledge sharing—helping others grow their skills
- **Influence research priorities and team strategy** while mentoring through technical leadership—establishing patterns, stakeholder coordination, reviewing code, and sharing expertise

You may be a good fit if you have:

- **Strong full-stack engineering experience**

- 6+ years building and shipping web products, with deep expertise across the stack
- Strong front-end skills: TypeScript/JavaScript, React, HTML/CSS, with an eye for clean, intuitive interface design
- Solid back-end and data pipeline experience: Python, API design, data analysis infrastructure
- Track record of independently driving features from prototype to production quality
- Comfort working as a technical "one-person shop" with full ownership

- **Deep conviction about human capability and agency**

- Strong perspective on how technology should enhance human capabilities rather than diminish them
- Experience or genuine passion for skill development, human-computer interaction, developer tools, or products that help people become more capable
- Understanding that skill development isn't tutorials or onboarding flows—it's experiences that make users progressively more capable, curious, and empowered over time
- Interest in how people develop mastery of AI itself and how that translates to broader capability growth
- Skepticism of purely engagement-driven metrics; interest in measuring skill development and capability outcomes

- **Research mindset with product execution skills**

- Comfort with ambiguity and exploring undefined problem spaces—you define what good looks like
- Experience working in research contexts or highly exploratory product environments
- Ability to rapidly prototype new ideas, test them with users, and iterate toward production
- Strong instincts for product design and user experience, even without formal design training

- Excitement about frontier thinking and new paradigms, not just optimization

- **Strategic thinking and focus**

- Experience setting vision for product projects, defining roadmaps with clear milestones, and maintaining team focus in exploratory environments
- Ability to balance research exploration with concrete deliverables and timelines
- Strong sense of prioritization—knowing what to build now, what to defer, and what to cut
- Track record of translating ambiguous opportunities into actionable plans

- **Technical leadership and communication**

- Experience providing technical direction, particularly in small or early-stage teams
- Ability to work effectively alongside researchers, translating between research insights and product implementation
- Strong written and verbal communication—you make complex technical and conceptual ideas accessible
- Track record of elevating team output through code quality, documentation, and knowledge sharing

Strong candidates may also have:

- Experience in developer tools, creative tools, learning platforms, or other products where user skill development and mastery matter more than time-on-site
- Background in learning sciences, cognitive science, HCI, skill acquisition research, or educational psychology (formal or self-directed)
- Experience with experimentation frameworks, A/B testing, or analytics that measure capability development in production
- Previous experience in research labs, frontier tech companies, or startups with high autonomy and ambiguity
- Published writing, talks, or open-source work on skill development, human-AI interaction, or product philosophy

- Experience building AI-native product experiences or working with LLMs in production contexts

What this role is not

This is a hands-on technical role building product features that ship to users, embedded within a research team. While you'll provide technical guidance and help set direction, this role doesn't involve people management or traditional team leadership responsibilities. If you're looking to transition into engineering management or lead a large team, this likely isn't the right fit.

The expected base compensation for this position is below. Our total compensation package for full-time employees includes equity, benefits, and may include incentive compensation.

Annual Salary:

\$320,000 - \$405,000 USD

Logistics

Education requirements: We require at least a Bachelor's degree in a related field or equivalent experience.

Location-based hybrid policy: Currently, we expect all staff to be in one of our offices at least 25% of the time. However, some roles may require more time in our offices.

Visa sponsorship: We do sponsor visas! However, we aren't able to successfully sponsor visas for every role and every candidate. But if we make you an offer, we will make every reasonable effort to get you a visa, and we retain an immigration lawyer to help with this.

We encourage you to apply even if you do not believe you meet every single qualification.

Not all strong candidates will meet every single qualification as listed. Research shows that people who identify as being from underrepresented groups are more prone to experiencing imposter syndrome and doubting the strength of their candidacy, so we urge you not to exclude yourself prematurely and to submit an application if you're interested in this work. We think AI systems like the ones we're building have enormous social and ethical implications. We think this makes representation even more important, and we strive to include a range of diverse perspectives on our team.

How we're different

We believe that the highest-impact AI research will be big science. At Anthropic we work as a single cohesive team on just a few large-scale research efforts. And we value impact — advancing our long-term goals of steerable, trustworthy AI — rather than work on smaller and more specific puzzles. We view AI research as an empirical science, which has as much in common with physics and biology as with traditional efforts in computer science. We're an extremely collaborative group, and we host frequent research discussions to ensure that we are pursuing the highest-impact work at any given time. As such, we greatly value communication skills.

The easiest way to understand our research directions is to read our recent research. This research continues many of the directions our team worked on prior to Anthropic, including: GPT-3, Circuit-Based Interpretability, Multimodal Neurons, Scaling Laws, AI & Compute, Concrete Problems in AI Safety, and Learning from Human Preferences.

Come work with us!

Anthropic is a public benefit corporation headquartered in San Francisco. We offer competitive compensation and benefits, optional equity donation matching, generous vacation and parental leave, flexible working hours, and a lovely office space in which to collaborate with colleagues. **Guidance on Candidates' AI Usage:** Learn about [our policy](#) for using AI in our application process



Create a Job Alert

Interested in building your career at Anthropic? Get future opportunities sent straight to your email.

[Create alert](#)

Apply for this job

* indicates a required field

Autofill with MyGreenhouse

First Name *

Last Name ***Email *****Country****Phone****Resume/CV**

Attach

Dropbox

Google Drive

Enter manually

Accepted file types: pdf, doc, docx, txt, rtf

(Optional) Personal Preferences

How do you pronounce your name?

Website**Are you open to working in-person in one of our offices 25% of the time? ***

Select...

When is the earliest you would want to start working with us?**Do you have any deadlines or timeline considerations we should be aware of?****AI Policy for Application ***

Select...

We believe that AI will have a transformative impact on the world, and we're seeking exceptional candidates who collaborate thoughtfully with Claude to realize this vision. At the same time, we want to understand your unique skills, expertise, and perspective through our hiring process. We invite you to review our [AI partnership guidelines for candidates](#) and confirm your understanding by selecting "Yes."

Why Anthropic? *

Why do you want to work at Anthropic? (We value this response highly - great answers are often 200-400 words.)

Do you require visa sponsorship? *

Select...

Will you now or will you in the future require employment visa sponsorship to work in the country in which the job you're applying for is located? *

Select...

Additional Information

Add a cover letter or anything else you want to share.

LinkedIn Profile

Please ensure to provide either your LinkedIn profile or Resume, we require at least one of the two.

Are you open to relocation for this role? *

Select...

What is the address from which you plan on working? If you would need to relocate, please type "relocating".

Have you ever interviewed at Anthropic before? *

Select... ▾

Voluntary Self-Identification

For government reporting purposes, we ask candidates to respond to the below self-identification survey. Completion of the form is entirely voluntary. Whatever your decision, it will not be considered in the hiring process or thereafter. Any information that you do provide will be recorded and maintained in a confidential file.

As set forth in Anthropic's Equal Employment Opportunity policy, we do not discriminate on the basis of any protected group status under any applicable law.

Gender

Select... ▾

Are you Hispanic/Latino?

Select... ▾

Race & Ethnicity Definitions

If you believe you belong to any of the categories of protected veterans listed below, please indicate by making the appropriate selection. As a government contractor subject to the Vietnam Era Veterans Readjustment Assistance Act (VEVRAA), we request this information in order to measure the effectiveness of the outreach and positive recruitment efforts we undertake pursuant to VEVRAA. Classification of protected categories is as follows:

A "disabled veteran" is one of the following: a veteran of the U.S. military, ground, naval or air service who is entitled to compensation (or who but for the receipt of military retired pay would be entitled to compensation) under laws administered by the Secretary of Veterans Affairs; or a person who was discharged or released from active duty because of a service-connected disability.

A "recently separated veteran" means any veteran during the three-year period beginning on the date of such veteran's discharge or release from active duty in the U.S. military, ground, naval, or air service.

An "active duty wartime or campaign badge veteran" means a veteran who served on active duty in the U.S. military, ground, naval or air service during a war, or in a campaign or

expedition for which a campaign badge has been authorized under the laws administered by the Department of Defense.

An "Armed forces service medal veteran" means a veteran who, while serving on active duty in the U.S. military, ground, naval or air service, participated in a United States military operation for which an Armed Forces service medal was awarded pursuant to Executive Order 12985.

Veteran Status

Select... ▼

Voluntary Self-Identification of Disability

Form CC-305

OMB Control Number 1250-0005

Page 1 of 1

Expires 04/30/2026

Why are you being asked to complete this form?

We are a federal contractor or subcontractor. The law requires us to provide equal employment opportunity to qualified people with disabilities. We have a goal of having at least 7% of our workers as people with disabilities. The law says we must measure our progress towards this goal. To do this, we must ask applicants and employees if they have a disability or have ever had one. People can become disabled, so we need to ask this question at least every five years.

Completing this form is voluntary, and we hope that you will choose to do so. Your answer is confidential. No one who makes hiring decisions will see it. Your decision to complete the form and your answer will not harm you in any way. If you want to learn more about the law or this form, visit the U.S. Department of Labor's Office of Federal Contract Compliance Programs (OFCCP) website at www.dol.gov/ofccp.

How do you know if you have a disability?

A disability is a condition that substantially limits one or more of your "major life activities." If you have or have ever had such a condition, you are a person with a disability. **Disabilities include, but are not limited to:**

- Alcohol or other substance use disorder (not currently using drugs illegally)
- Autoimmune disorder, for example, lupus, fibromyalgia, rheumatoid arthritis, HIV/AIDS
- Blind or low vision
- Cancer (past or present)

- Cardiovascular or heart disease
- Celiac disease
- Cerebral palsy
- Deaf or serious difficulty hearing
- Diabetes
- Disfigurement, for example, disfigurement caused by burns, wounds, accidents, or congenital disorders
- Epilepsy or other seizure disorder
- Gastrointestinal disorders, for example, Crohn's Disease, irritable bowel syndrome
- Intellectual or developmental disability
- Mental health conditions, for example, depression, bipolar disorder, anxiety disorder, schizophrenia, PTSD
- Missing limbs or partially missing limbs
- Mobility impairment, benefiting from the use of a wheelchair, scooter, walker, leg brace(s) and/or other supports
- Nervous system condition, for example, migraine headaches, Parkinson's disease, multiple sclerosis (MS)
- Neurodivergence, for example, attention-deficit/hyperactivity disorder (ADHD), autism spectrum disorder, dyslexia, dyspraxia, other learning disabilities
- Partial or complete paralysis (any cause)
- Pulmonary or respiratory conditions, for example, tuberculosis, asthma, emphysema
- Short stature (dwarfism)
- Traumatic brain injury

Disability StatusSelect... 

PUBLIC BURDEN STATEMENT: According to the Paperwork Reduction Act of 1995 no persons are required to respond to a collection of information unless such collection displays a valid OMB control number. This survey should take about 5 minutes to complete.

[Submit application](#)