



Anthropic

California, United States (Hybrid)

Apply Now

Senior Human-AI Collaboration Research Engineer, Education Labs

Anthropic • California, United States • Hybrid • 3 Weeks ago • 6 Years + • \$320,000 PA - \$405,000 PA

Research Development

Job Summary



Job Description



Anthropic's Education Labs team is seeking a Senior Human-AI Collaboration Research Engineer to explore how AI transforms human capability development and build new product paradigms. This role involves operating with research-level autonomy while shipping features to millions of users, sitting at the intersection of research and product execution. The engineer will define technical direction, prototype new ideas, and build experiences that make users more capable and thoughtful, focusing on skill development with AI.

Good To Have:

- Experience in developer tools, creative tools, or learning platforms.
- Background in learning sciences, cognitive science, HCI, or educational psychology.
- Experience with experimentation frameworks, A/B testing, or analytics for capability development.
- Previous experience in research labs, frontier tech companies, or startups.
- Published writing, talks, or open-source work on skill development or human-AI interaction.
- Experience building AI-native product experiences or working with LLMs.

Must Have:

- Build and ship features for user skill development with AI.
- Architect and implement full-stack prototypes (front-end and back-end).
- Define technical direction and quality standards for product work.

Apply Now

- Document and share work through technical writing, prototypes, and presentations.
- Elevate team craft through thoughtful code reviews and knowledge sharing.

- Influence research priorities and team strategy through technical leadership.
- 6+ years of full-stack engineering experience.
- Strong front-end skills: TypeScript/JavaScript, React, HTML/CSS.
- Solid back-end and data pipeline experience: Python, API design, data analysis infrastructure.
- Track record of independently driving features from prototype to production quality.
- Strong product instincts and clean interface design sensibilities.
- Comfort with ambiguity and exploring undefined problem spaces.
- Ability to rapidly prototype new ideas, test with users, and iterate.
- Experience setting vision for product projects and defining roadmaps.
- Strong written and verbal communication skills.

Perks:

- Competitive compensation
- Benefits
- Optional equity donation matching
- Generous vacation
- Parental leave
- Flexible working hours
- Lovely office space for collaboration

17 skills required for this role

Add these skills to join the top 1% applicants for this job

Team Management + Communication + Data Analytics + Unity +

Talent Acquisition + game texts + Alpha/Beta Testing + React + HTML +

User Experience (UX) + Prototyping + Front-End + Back-End + Python +

TypeScript + JavaScript + CSS +

Job Details



About Anthropic

Anthropic's mission is to create reliable, interpretable, and steerable AI systems. We want AI to be safe and beneficial for our users and for society as a whole. Our team is a quickly growing group of committed researchers, engineers, policy experts, and business leaders working together to build beneficial AI systems.

About the role

Education Labs is a research-driven team with a product mandate: we study how AI transforms human capability development, and we build new product paradigms based on what we discover. As a Senior Product Research Engineer, you'll be the second technical builder on this small, mission-driven team—operating with research-level autonomy while shipping features that reach millions of users.

This role sits at the intersection of research exploration and product execution. You'll work alongside researchers studying the societal impacts of AI, while building and shipping features that help people develop skills—both in working with AI itself and in their broader capabilities. You'll have significant creative license to define what "good" looks like, exploring new interaction paradigms rather than optimizing existing patterns.

Our team believes that developing skill with AI is fundamental to human agency and flourishing in an AI-enabled world. We're building experiences that make users more capable, more thoughtful, and more empowered—not just more engaged. This means integrating skill development principles into product design, leveraging Claude itself as part of the capability-building journey, and measuring success by how users grow in their ability to work with AI and apply those skills to their own goals.

You'll be something of a one-person technical shop with full-stack ownership, prototyping new ideas, establishing technical direction, and shipping production-quality features. You'll need strong product instincts and clean interface design sensibilities, balanced with comfort in ambiguity and frontier thinking.

Responsibilities:

- **Build and ship features** that help users develop skills in working with AI and translate those skills into their own capability growth—measuring success by skill development and product mastery, not just engagement
- **Architect and implement end-to-end prototypes** (front-end and back-end) that test new interaction paradigms for skill development with AI, including novel uses of Claude as a skill-building partner
- **Collaborate with researchers** to translate insights about skill development and human-AI collaboration into shipped product experiences

- **Work closely with product, design, and research teams** across Anthropic to influence how skill development principles shape our broader product strategy
- **Analyze product data and user behavior** using quantitative and qualitative methods to understand how users develop skills with AI and what actually helps them become more capable
- **Document and share your work** through clear technical writing, prototypes, and presentations that influence thinking across the organization
- **Elevate team craft** through thoughtful code reviews, technical discussions, and knowledge sharing—helping others grow their skills
- **Influence research priorities and team strategy** while mentoring through technical leadership—establishing patterns, stakeholder coordination, reviewing code, and sharing expertise

You may be a good fit if you have:

- **Strong full-stack engineering experience**
 - 6+ years building and shipping web products, with deep expertise across the stack
 - Strong front-end skills: TypeScript/JavaScript, React, HTML/CSS, with an eye for clean, intuitive interface design
 - Solid back-end and data pipeline experience: Python, API design, data analysis infrastructure
 - Track record of independently driving features from prototype to production quality
 - Comfort working as a technical "one-person shop" with full ownership
- **Deep conviction about human capability and agency**
 - Strong perspective on how technology should enhance human capabilities rather than diminish them
 - Experience or genuine passion for skill development, human-computer interaction, developer tools, or products that help people become more capable
 - Understanding that skill development isn't tutorials or onboarding flows—it's experiences that make users progressively more capable, curious, and empowered over time
 - Interest in how people develop mastery of AI itself and how that translates to broader capability growth
 - Skepticism of purely engagement-driven metrics; interest in measuring skill development and capability outcomes
- **Research mindset with product execution skills**
 - Comfort with ambiguity and exploring undefined problem spaces—you define what good looks like
 - ~~Experience working in research contexts or highly exploratory product environments~~
 - Strong instincts for product design and user experience, even without formal design training
 - Excitement about frontier thinking and new paradigms, not just optimization

- **Strategic thinking and focus**

- Experience setting vision for product projects, defining roadmaps with clear milestones, and maintaining team focus in exploratory environments
- Ability to balance research exploration with concrete deliverables and timelines
- Strong sense of prioritization—knowing what to build now, what to defer, and what to cut
- Track record of translating ambiguous opportunities into actionable plans

- **Technical leadership and communication**

- Experience providing technical direction, particularly in small or early-stage teams
- Ability to work effectively alongside researchers, translating between research insights and product implementation
- Strong written and verbal communication—you make complex technical and conceptual ideas accessible
- Track record of elevating team output through code quality, documentation, and knowledge sharing

Strong candidates may also have:

- Experience in developer tools, creative tools, learning platforms, or other products where user skill development and mastery matter more than time-on-site
- Background in learning sciences, cognitive science, HCI, skill acquisition research, or educational psychology (formal or self-directed)
- Experience with experimentation frameworks, A/B testing, or analytics that measure capability development in production
- Previous experience in research labs, frontier tech companies, or startups with high autonomy and ambiguity
- Published writing, talks, or open-source work on skill development, human-AI interaction, or product philosophy
- Experience building AI-native product experiences or working with LLMs in production contexts

What this role is not

This is a hands-on technical role building product features that ship to users, embedded within a research team. While you'll provide technical guidance and help set direction, this role doesn't involve people management or traditional team leadership responsibilities. If you're looking to transition into engineering management or lead a large team, this likely isn't the right fit.

Annual Salary:

\$320,000 - \$405,000 USD

Logistics

Education requirements: We require at least a Bachelor's degree in a related field or equivalent experience.

Location-based hybrid policy: Currently, we expect all staff to be in one of our offices at least 25% of the time. However, some roles may require more time in our offices.

Visa sponsorship: We do sponsor visas! However, we aren't able to successfully sponsor visas for every role and every candidate. But if we make you an offer, we will make every reasonable effort to get you a visa, and we retain an immigration lawyer to help with this.

We encourage you to apply even if you do not believe you meet every single qualification. Not all strong candidates will meet every single qualification as listed. Research shows that people who identify as being from underrepresented groups are more prone to experiencing imposter syndrome and doubting the strength of their candidacy, so we urge you not to exclude yourself prematurely and to submit an application if you're interested in this work. We think AI systems like the ones we're building have enormous social and ethical implications. We think this makes representation even more important, and we strive to include a range of diverse perspectives on our team.

How we're different

We believe that the highest-impact AI research will be big science. At Anthropic we work as a single cohesive team on just a few large-scale research efforts. And we value impact — advancing our long-term goals of steerable, trustworthy AI — rather than work on smaller and more specific puzzles. We view AI research as an empirical science, which has as much in common with physics and biology as with traditional efforts in computer science. We're an extremely collaborative group, and we host frequent research discussions to ensure that we are pursuing the highest-impact work at any given time. As such, we greatly value communication skills.

The easiest way to understand our research directions is to read our recent research. This research continues many of the directions our team worked on prior to Anthropic, including: GPT-3, Circuit-Based Interpretability, Multimodal Neurons, Scaling Laws, AI & Compute, Concrete Problems in AI Safety, and Learning from Human Preferences.

Come work with us!

Anthropic is a public benefit corporation headquartered in San Francisco. We offer competitive compensation

about our policy for using AI in our application process

Job Alerts



Set alerts for more jobs like **Senior Human-AI Collaboration Research Engineer, Education Labs**



Set alerts for new jobs by **Anthropic**



Set alerts for new **Research Development jobs in United States**



Set alerts for new jobs in **United States**



Set alerts for **Research Development (Remote)** jobs



Manage Alerts

Similar Jobs



More Research Development Jobs

Senior Responsible AI Research Scientist....
Autodesk
Portland, Oregon, United States...
🕒 25 days ago

AI Enablement Lead, People and Places
Autodesk
Denver, Colorado, United State...
🕒 25 days ago

II A P C
🕒 25 days ago

More Software Development & Engineering Jobs

Staff Software Engineer
nubank
Mexico City, Mexico (Hybrid)
🕒 25 days ago

Staff Software Engineer
nubank
Buenos Aires, Ciudad Autónom...
🕒 25 days ago

S (J M C
🕒 25 days ago



Anthropic

51 Active Jobs

Creative Director, Enterprise

San Francisco, California, United States (hybrid)

26 days ago

Machine Learning Systems Engineer, Research Tools

United States (hybrid)

27 days ago

Machine Learning Systems Engineer, RL Engineering

San Francisco, California, United States (hybrid)

27 days ago

Software Engineer, Biology & Life Sciences

San Francisco, California, United States (hybrid)

[View All Jobs](#) 

Level Up Your Career in Game Development!

Transform Your Passion into Profession with Our Comprehensive Courses for Aspiring Game Developers.



[Upskill Now](#)

Company

[About Us](#)

[Terms Of Service](#)

Contact Us

 hello@outscal.com



[Refund Policy](#)

[Pricing Policy](#)

Cancellation Policy
